

Spotify Classification

February 21, 2021

0.1 Description of the final model

The final model for the classification problem uses an ensemble learning method. This combines different predictors that were explored in this report. This allows for better predictions than with a single predictor that does pretty well. To measure the performance of a predictor, precision, recall, and f1 scores are used. The voting classifier has a precision of 14%, recall of 25% and f1 score of 17%. The voting classifier was used as the 3 different methods used do not classify the genre well as they make different errors. The classifier uses hard voting which predicts the class that gets the most votes.

0.1.1 Importing packages and datasets

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: Test = pd.read_csv('CS98XClassificationTest.csv')
Train = pd.read_csv('CS98XClassificationTrain.csv')
```

0.1.2 Viewing data

```
[3]: Test.head(2)
```

```
[3]:
```

	Id		title \
0	454		Pump It
1	455		Circle of Life - From "The Lion King"/Soundtra...

		artist	year	bpm	nrgy	dnce	dB	live	val	dur	acous \
0		The Black Eyed Peas	2005	154	93	65	-3	75	74	213	1
1		Elton John	1994	161	39	30	-15	11	14	292	26

	spch	pop
0	18	72
1	3	59

```
[4]: Train.head(2)
```

```
[4]: Id          title          artist  year  bpm  nrgy  dnce  dB  live  \
0    1      My Happiness  Connie Francis  1996  107   31   45  -8   13
1    2  Unchained Melody  The Teddy Bears  2011  114   44   53  -8   13

      val  dur  acous  spch  pop      top genre
0    28  150    75    3   44  adult standards
1    47  139    49    3   37                NaN
```

```
[5]: Test.shape
```

```
[5]: (113, 14)
```

```
[6]: Train.shape
```

```
[6]: (453, 15)
```

```
[7]: Train.describe()
```

```
[7]:
```

	Id	year	bpm	nrgy	dnce	\
count	453.000000	453.000000	453.000000	453.000000	453.000000	
mean	227.000000	1991.443709	118.399558	60.070640	59.565121	
std	130.914094	16.776103	25.238713	22.205284	15.484458	
min	1.000000	1948.000000	62.000000	7.000000	18.000000	
25%	114.000000	1976.000000	100.000000	43.000000	49.000000	
50%	227.000000	1994.000000	119.000000	63.000000	61.000000	
75%	340.000000	2007.000000	133.000000	78.000000	70.000000	
max	453.000000	2019.000000	199.000000	100.000000	96.000000	

	dB	live	val	dur	acous	spch	\
count	453.000000	453.000000	453.000000	453.000000	453.000000	453.000000	
mean	-8.836645	17.757174	59.465784	226.278146	32.982340	5.660044	
std	3.577187	13.830300	24.539868	63.770380	29.530015	5.550581	
min	-24.000000	2.000000	6.000000	98.000000	0.000000	2.000000	
25%	-11.000000	9.000000	42.000000	181.000000	7.000000	3.000000	
50%	-8.000000	13.000000	61.000000	223.000000	24.000000	4.000000	
75%	-6.000000	23.000000	80.000000	262.000000	58.000000	6.000000	
max	-1.000000	93.000000	99.000000	511.000000	100.000000	47.000000	

	pop
count	453.000000
mean	60.743929
std	13.470083
min	26.000000
25%	53.000000
50%	63.000000
75%	71.000000
max	84.000000


```
[ ]: Train.dtypes
```

```
[ ]: Train.isnull().sum()
```

```
[ ]: Test.isnull().sum()
```

```
[56]: plt.figure(figsize =(10,4))  
correlation = Train.corr()  
sns.heatmap(correlation, annot=True)
```

```
[56]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe5ee862ee0>
```



1 Preparing the data

To get the data tables ready for the analysis, there are some columns that will not work with the methods as they prefer integers. The Id, title, artist, and year are variables that will not add to the classification analysis as they would not impact the classification of a song.

There are also 15 missing values from the top genre column in the training dataset. Since the missing values are less than 20%, this will not affect the analysis too much so will be dropped to make for improved classification.

```
[12]: Train = Train.dropna(subset=["top genre"])
```

```
[13]: Train.head(2)
```

```
[13]:   Id      title      artist  year  bpm  nrgy  dnce  dB  live  \  
0    1  My Happiness  Connie Francis  1996  107   31   45  -8   13  
2    3  How Deep Is Your Love      Bee Gees  1979  105   36   63  -9   13
```

	val	dur	acous	spch	pop	top genre
0	28	150	75	3	44	adult standards
2	67	245	11	3	77	adult standards

```
[14]: spotify_cat = Train[["top genre"]]
      spotify_cat.head(2)
```

```
[14]:          top genre
0  adult standards
2  adult standards
```

```
[15]: from sklearn.preprocessing import OrdinalEncoder
      ordinal_encoder = OrdinalEncoder()
      spotify_cat_encoded = ordinal_encoder.fit_transform(spotify_cat)
      spotify_cat_encoded[:5]
```

```
[15]: array([[ 1.],
             [ 1.],
             [ 1.],
             [68.],
             [80.]])
```

```
[16]: from sklearn.preprocessing import OneHotEncoder
      spotify_encoder = OneHotEncoder()
      spotify_cat_1hot = spotify_encoder.fit_transform(spotify_cat)
      spotify_cat_1hot
```

```
[16]: <438x86 sparse matrix of type '<class 'numpy.float64'>'
      with 438 stored elements in Compressed Sparse Row format>
```

```
[ ]: spotify_cat_1hot.toarray()
```

```
[18]: y= spotify_cat_encoded
```

```
[19]: X = Train.drop(['title', 'artist', 'Id', 'year', 'top genre', 'pop'], axis=1)
```

```
[57]: X.head(2)
```

```
[57]:    bpm  nrgy  dnce  dB  live  val  dur  acous  spch
0   107   31   45  -8   13   28  150   75   3
2   105   36   63  -9   13   67  245   11   3
```

1.1 Scaling and splitting data

Machine learning techniques prefer to work on numerical attributes with similar scales. Standardisation is less affected by outliers in the data.

```
[21]: from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import StandardScaler
      from sklearn.impute import SimpleImputer
```

```
[22]: from sklearn.model_selection import train_test_split

      X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,
      ↪random_state=42)
```

```
[23]: scaler = StandardScaler()
      scaler.fit(X_train)
      X_train = scaler.transform(X_train)
      X_test = scaler.transform(X_test)
```

```
[ ]: X_train
```

```
[ ]: X_test
```

2 Select and training models

2.1 Using logistic regression

```
[ ]: from sklearn.linear_model import LogisticRegression

      log_reg = LogisticRegression()
      log_reg.fit(X_train,y_train)
```

```
[27]: pred = log_reg.predict(X_test)
```

```
[ ]: pred
```

2.1.1 Evaluation of logistic regression

```
[29]: from sklearn.metrics import confusion_matrix
      conf = confusion_matrix(y_test, pred)
```

```
[30]: conf
```

```
[30]: array([[10,  0,  3, ...,  0,  0,  0],
           [ 0,  0,  0, ...,  0,  0,  0],
           [ 0,  0,  6, ...,  0,  0,  0],
           ...,
           [ 0,  0,  1, ...,  0,  0,  0],
           [ 1,  0,  0, ...,  0,  0,  0],
           [ 1,  0,  0, ...,  0,  0,  0]])
```

We are able to produce a confusion matrix where each row represents an actual class and the column shows predicted class. The full confusion matrix can't be viewed so we will use a classification report to get a measure of how well the model performs.

```
[31]: from sklearn.metrics import classification_report
```

```
[32]: class_rep = classification_report(y_test, pred)
      print("Logistic Regression: \n", class_rep)
```

Logistic Regression:

	precision	recall	f1-score	support
1.0	0.50	0.62	0.56	16
3.0	0.00	0.00	0.00	0
4.0	0.19	0.60	0.29	10
6.0	0.00	0.00	0.00	1
9.0	0.00	0.00	0.00	3
10.0	0.00	0.00	0.00	2
13.0	0.00	0.00	0.00	1
14.0	0.00	0.00	0.00	1
16.0	0.00	0.00	0.00	2
18.0	0.00	0.00	0.00	1
20.0	0.00	0.00	0.00	1
24.0	0.00	0.00	0.00	3
25.0	0.00	0.00	0.00	3
31.0	0.00	0.00	0.00	1
32.0	0.00	0.00	0.00	1
36.0	0.00	0.00	0.00	0
39.0	0.00	0.00	0.00	1
40.0	0.00	0.00	0.00	1
42.0	0.00	0.00	0.00	1
43.0	0.00	0.00	0.00	1
46.0	0.00	0.00	0.00	1
47.0	0.00	0.00	0.00	1
49.0	0.00	0.00	0.00	1
51.0	0.42	0.77	0.54	13
52.0	0.00	0.00	0.00	3
53.0	0.00	0.00	0.00	1
56.0	0.00	0.00	0.00	1
57.0	0.00	0.00	0.00	1
60.0	0.50	1.00	0.67	1
61.0	0.00	0.00	0.00	1
62.0	0.00	0.00	0.00	3
63.0	0.00	0.00	0.00	2
68.0	0.00	0.00	0.00	6
70.0	0.00	0.00	0.00	0
80.0	0.00	0.00	0.00	1
83.0	0.00	0.00	0.00	1

	85.0	0.00	0.00	0.00	1
accuracy				0.31	88
macro avg	0.04	0.08	0.06		88
weighted avg	0.18	0.31	0.22		88

```
/Users/Evelyn/opt/anaconda3/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/Users/Evelyn/opt/anaconda3/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in labels with no true samples.
Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

From the classification report, different values are obtained about the models performance. The overall precision is 0.18, recall 0.27, and f1 score is 0.21. The f1 score is low because both recall and precision scores are also low.

Logistic regression will be used as a baseline to compare more complex methods for classification.

2.2 Decision tree classifier

Decision tree classification will be used in this section to see if there is an improvement to the logistic regression model.

```
[33]: from sklearn.tree import DecisionTreeClassifier
```

With the remaining model methods, the target variable will not be changed to ensure no information is lost. The methods are also able to handle categorical variables.

```
[34]: X_train, X_test, y_train, y_test = train_test_split(X,spotify_cat, test_size=0.
    ↪2, random_state=42)
```

```
[35]: tree = DecisionTreeClassifier()
tree.fit(X_train,y_train)

pred_tree = tree.predict(X_test)
```

```
[ ]: pred_tree
```

The model seems to be good at predicting songs that belong to rock and pop sections.

2.2.1 Evaluation of decision tree

```
[37]: confusion_matrix(y_test, pred_tree)
```

```
[37]: array([[6, 0, 2, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          [1, 0, 4, ..., 0, 0, 0],
          ...,
          [0, 0, 1, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0]])
```

```
[38]: class_rep1 = classification_report(y_test, pred_tree)
      print("Decision Tree: \n", class_rep1)
```

Decision Tree:

	precision	recall	f1-score	support
adult standards	0.40	0.38	0.39	16
afropop	0.00	0.00	0.00	0
album rock	0.24	0.40	0.30	10
alternative rock	0.00	0.00	0.00	1
american folk revival	0.00	0.00	0.00	0
art pop	0.00	0.00	0.00	0
art rock	0.00	0.00	0.00	3
atl hip hop	0.00	0.00	0.00	2
avant-garde jazz	0.00	0.00	0.00	1
barbadian pop	0.00	0.00	0.00	1
bebop	0.00	0.00	0.00	2
belgian pop	0.00	0.00	0.00	1
blues	0.00	0.00	0.00	1
boogaloo	0.00	0.00	0.00	0
boy band	0.12	0.33	0.18	3
brill building pop	0.00	0.00	0.00	3
british invasion	0.00	0.00	0.00	1
british soul	0.00	0.00	0.00	1
canadian folk	0.00	0.00	0.00	0
canadian pop	0.00	0.00	0.00	1
celtic rock	0.00	0.00	0.00	1
chanson	0.00	0.00	0.00	0
chicago rap	0.00	0.00	0.00	1
chicago soul	0.00	0.00	0.00	1
classic rock	0.00	0.00	0.00	1
classic soul	0.00	0.00	0.00	1
classic uk pop	0.00	0.00	0.00	0
country	0.00	0.00	0.00	1
dance pop	0.09	0.08	0.08	13
dance rock	0.00	0.00	0.00	3

deep adult standards	1.00	1.00	1.00	1
deep house	0.00	0.00	0.00	0
disco	0.00	0.00	0.00	1
disco house	0.00	0.00	0.00	1
east coast hip hop	1.00	1.00	1.00	1
eurodance	0.00	0.00	0.00	1
europop	0.00	0.00	0.00	3
g funk	0.00	0.00	0.00	2
glam rock	0.00	0.00	0.00	6
latin	0.00	0.00	0.00	0
mellow gold	0.00	0.00	0.00	0
new wave pop	0.00	0.00	0.00	0
permanent wave	0.00	0.00	0.00	0
pop	0.00	0.00	0.00	1
soft rock	0.00	0.00	0.00	1
yodeling	0.00	0.00	0.00	1
accuracy			0.16	88
macro avg	0.06	0.07	0.06	88
weighted avg	0.14	0.16	0.15	88

```

/Users/Evelyn/opt/anaconda3/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/Users/Evelyn/opt/anaconda3/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in labels with no true samples.
Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

The classification report tells us how well the decision tree method performs. With a precision of 14%, recall of 15% and f1 score of 14%.

2.3 Random forest

Random forest are able to handle categorical values so will be used on the data to build a new model. There are 2 candidate models but random forest is an ensemble method to get a better predictor.

```
[39]: from sklearn.ensemble import RandomForestClassifier
```

```
[40]: forest = RandomForestClassifier(
        min_samples_leaf=40, criterion = 'entropy',
        n_estimators=15,
        max_depth = 7,
        oob_score=True,
```

```

n_jobs=-1,
random_state=42,
max_features='auto')

```

```
[ ]: forest.fit(X_train, y_train)
```

```
[42]: forest_tree = forest.predict(X_test)
```

```
[ ]: forest_tree
```

2.3.1 Evaluation of random forest

```

[44]: class_rep2 = classification_report(y_test, forest_tree)
print("Random Forest: \n", class_rep2)

```

Random Forest:

	precision	recall	f1-score	support
adult standards	0.37	0.69	0.48	16
album rock	0.18	0.70	0.29	10
alternative rock	0.00	0.00	0.00	1
art rock	0.00	0.00	0.00	3
atl hip hop	0.00	0.00	0.00	2
avant-garde jazz	0.00	0.00	0.00	1
barbadian pop	0.00	0.00	0.00	1
bebop	0.00	0.00	0.00	2
belgian pop	0.00	0.00	0.00	1
blues	0.00	0.00	0.00	1
boy band	0.00	0.00	0.00	3
brill building pop	0.00	0.00	0.00	3
british invasion	0.00	0.00	0.00	1
british soul	0.00	0.00	0.00	1
canadian pop	0.00	0.00	0.00	1
celtic rock	0.00	0.00	0.00	1
chicago rap	0.00	0.00	0.00	1
chicago soul	0.00	0.00	0.00	1
classic rock	0.00	0.00	0.00	1
classic soul	0.00	0.00	0.00	1
country	0.00	0.00	0.00	1
dance pop	0.21	0.31	0.25	13
dance rock	0.00	0.00	0.00	3
deep adult standards	0.00	0.00	0.00	1
disco	0.00	0.00	0.00	1
disco house	0.00	0.00	0.00	1
east coast hip hop	0.00	0.00	0.00	1
eurodance	0.00	0.00	0.00	1
europop	0.00	0.00	0.00	3

g funk	0.00	0.00	0.00	2
glam rock	0.00	0.00	0.00	6
pop	0.00	0.00	0.00	1
soft rock	0.00	0.00	0.00	1
yodeling	0.00	0.00	0.00	1
accuracy			0.25	88
macro avg	0.02	0.05	0.03	88
weighted avg	0.12	0.25	0.16	88

```
/Users/Evelyn/opt/anaconda3/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

From the classification report, we are able to extract the precision, recall, and f1 score. Random forest method does moderately well with an f1 score of 16%. However, logistic regression still has the best score.

3 Using ensemble classifier to improve model

We will attempt to use a voting classifier. The predictors used so far have an okay level of accuracy but can be improved on by using an ensemble method. This will bring together the predictions of all classifiers and use the class that gets the most votes. In this problem, hard voting is used.

```
[ ]: from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.linear_model import LogisticRegression

log_reg = LogisticRegression()
rnd_clf = RandomForestClassifier()
tree = DecisionTreeClassifier()
voting_clf = VotingClassifier(
    estimators=[('lr', log_reg), ('rf', rnd_clf), ('dc', tree)], voting='hard')
voting_clf.fit(X_train, y_train)
```

```
[46]: predict = voting_clf.predict(X_test)
```

```
[ ]: predict
```

```
[48]: class_rep4 = classification_report(y_test, predict)
print("Voting Classifier: \n", class_rep4)
```

Voting Classifier:

precision	recall	f1-score	support
-----------	--------	----------	---------

adult standards	0.43	0.62	0.51	16
afropop	0.00	0.00	0.00	0
album rock	0.19	0.70	0.30	10
alternative rock	0.00	0.00	0.00	1
art rock	0.00	0.00	0.00	3
atl hip hop	0.00	0.00	0.00	2
avant-garde jazz	0.00	0.00	0.00	1
barbadian pop	0.00	0.00	0.00	1
bebop	0.00	0.00	0.00	2
belgian pop	0.00	0.00	0.00	1
blues	0.00	0.00	0.00	1
boy band	0.00	0.00	0.00	3
brill building pop	0.00	0.00	0.00	3
british invasion	0.00	0.00	0.00	1
british soul	1.00	1.00	1.00	1
bubblegum dance	0.00	0.00	0.00	0
canadian pop	0.00	0.00	0.00	1
celtic rock	0.00	0.00	0.00	1
chicago rap	0.00	0.00	0.00	1
chicago soul	0.00	0.00	0.00	1
classic rock	0.00	0.00	0.00	1
classic soul	0.00	0.00	0.00	1
country	0.00	0.00	0.00	1
dance pop	0.38	0.46	0.41	13
dance rock	0.00	0.00	0.00	3
deep adult standards	0.00	0.00	0.00	1
detroit hip hop	0.00	0.00	0.00	0
disco	0.00	0.00	0.00	1
disco house	0.00	0.00	0.00	1
east coast hip hop	1.00	1.00	1.00	1
eurodance	0.00	0.00	0.00	1
europop	0.00	0.00	0.00	3
g funk	0.00	0.00	0.00	2
glam rock	0.00	0.00	0.00	6
pop	0.00	0.00	0.00	1
soft rock	0.00	0.00	0.00	1
yodeling	0.00	0.00	0.00	1
accuracy			0.28	88
macro avg	0.08	0.10	0.09	88
weighted avg	0.18	0.28	0.21	88

```

/Users/Evelyn/opt/anaconda3/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

```

```
/Users/Evelyn/opt/anaconda3/lib/python3.8/site-  
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning: Recall  
and F-score are ill-defined and being set to 0.0 in labels with no true samples.  
Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))
```

The classification report from the voting classifier shows a good enough performance. The recall is highest at 25% and the precision is 14%. The f1 score is 17%.

4 Evaluation on test set

```
[49]: Test1 = Test.drop(['Id', 'title', 'artist', 'year', 'pop'], axis=1)
```

```
[50]: Test1.head(2)
```

```
[50]:
```

	bpm	nrgy	dnce	dB	live	val	dur	acous	spch
0	154	93	65	-3	75	74	213	1	18
1	161	39	30	-15	11	14	292	26	3

```
[51]: final_predictions = voting_clf.predict(Test1)
```

```
[ ]: final_predictions
```

5 Converting data to CSV

```
[53]: submission = pd.DataFrame({"Id":Test["Id"], "top genre": final_predictions})  
      submission.head(3)
```

```
[53]:
```

	Id	top genre
0	454	big room
1	455	album rock
2	456	adult standards

```
[54]: file = "ClassificationPredictions.csv"  
  
      submission.to_csv(file, index = False, header = 1)
```

6 Conclusion

The following model that has been chosen for the classification problem was the best out of all of the tested models/methods. To start with, at the beginning of this report, data cleaning was vital to the successful analysis. Certain columns were removed from the training dataset as they would not be of any use as these factors in hindsight do not contribute to the genre a song belongs in.

The different methods used in this report have all been able to build a successful ML algorithms. A voting classifier algorithm is the final model as it achieves the best recall, f1, and accuracy score. This is successful by combining predictions from different classifiers.

From the Kaggle InClass competition, the final model achieves a score of 0.28571 thus concluding the model built has been successful in predicting the classification of a song. This is a better score than on the training set.